

PROFESSIONAL SUMMARY

- 9+ years of experience designing, building, and scaling enterprise-grade data platforms on Azure, specializing in lakehouse architecture, distributed data processing, and cloud-native ETL/ELT workloads.
- Spearheaded the Azure Data Factory (ADF) ecosystem by architecting a config-driven, metadata-based ingestion framework to accelerate onboarding of new source systems and reduce development effort across teams.
- Designed and implemented end-to-end lakehouse solutions using ADF, Azure Databricks, and ADLS Gen2 to deliver scalable, modular, reusable data pipelines across business domains.
- Modernized legacy platforms by migrating a 4TB+ enterprise data warehouse to Azure Synapse Analytics (120+ pipelines/stored procedures; ~80 dimension/fact tables), ensuring minimal downtime and business continuity.
- Established engineering standards (CI/CD, branching, code reviews, and DevOps practices) to support reliable, parallel development across teams.
- Owned production stability and system reliability through monitoring, fault detection, and continuous improvements to meet SLAs.
- Partnered with InfoSec and Infrastructure to implement platform security, secure connectivity, and access controls.
- Managed and mentored junior data engineers, providing technical and operational guidance to support delivery.
- Conducted L1/L2 technical interviews for senior/junior data engineering roles, primarily focused on the Azure data platform.
- Served as the Databricks Partner Connect lead, strengthening the partnership and delivering targeted PoCs to demonstrate technical capability.

KEY SKILLS

Data Platform & Architecture

- Lakehouse Architecture
- Enterprise Data Warehousing
- Data Platform Design

Azure Data Ecosystem

- Azure Databricks
- Azure Data Factory
- Azure Synapse Analytics
- Azure SQL Database/Warehouse
- Azure Data Lake Storage Gen2

Data Engineering & Processing

- Apache Spark/PySpark
- Spark SQL & Delta Lake Architecture
- ETL & ELT Design Paradigm
- Data Quality/Validation Frameworks

DataOps, CI/CD & Infrastructure

- Azure DevOps
- IaC (ARM Template Deployment)
- Git based version Control

Programming & Querying

- Python
- T-SQL

Leadership & Delivery

- Data Engineering Team Leadership
- Stakeholder Management
- Solution Architecture
- Platform Migration
- Technical Hiring (L1&L2)

EDUCATION

B.Tech. (Hons) in Computer Science & Engineering (07/2013 – 06/2017)
JECRC Jaipur, Rajasthan Technical University

HSC (2013)
Central Board of Secondary Education

SSC (2011)
Central Board of Secondary Education

Position of Responsibilities

Jan'24 – Present	Data Engineering Manager, Data Management
Jan'22 - Dec'23	Senior Consultant, Data Management
Jun'20 - Dec'21	Consultant, Data Management

Lead Data Engineer & ETL Developer [Domain: Pharmaceuticals (US)]

- Lead ADF infrastructure for Lakehouse ingestion—building a config-driven framework and leading end-to-end development of 50+ pipelines.
- Connected 8+ source systems (private endpoints) and built 20+ parameterized pipelines for full/incremental loads.
- Standardized raw ingestion into Azure Data Lake Storage (ADLS) and transformed data into Delta tables using Azure Databricks.
- Built reusable trigger/orchestration pipelines to scale ingestion across plants, making new onboardings repeatable and low-effort.
- Set up Splunk-based monitoring and email alerting for trigger/activity failures, reducing time-to-detect and improving incident response.
- Developed T-SQL stored procedures in Azure SQL DB to capture ADF metadata and enable centralized, auditable logging.
- Built daily maintenance pipelines to optimize/vacuum Delta tables and clean up cloud storage (retention-based purges) to control storage and compute costs.
- Authored PySpark notebooks in Azure Databricks to process raw data and load curated Delta tables.
- Implemented Medallion architecture (Bronze/Silver/Gold) to structure the Lakehouse for scalable transformations and consumption.
- Designed an IaC-based CI/CD framework for ADF pipelines and datasets, enabling selective (partial) promotions across environments.

Sr Data Engineer & DevOps Engineer [Domain: Auto Auctions (US)]

- Led migration of a legacy on-prem BI platform to Azure for a major US auto-auctions firm.
- Re-engineered SSIS packages into Azure Synapse pipelines while preserving legacy behavior.
- Converted stored procedures to Azure SQL DW-compatible T-SQL and validated parity for all dimension and fact loads.
- Partnered with the platform owner to stabilize the cloud platform and ensure executive dashboards remained accurate and reliable.
- Troubleshoot and fixed 60+ stored procedures to keep KPI tables consistent across all dimension and fact sources.
- Built ingestion and validation pipelines to quarantine corrupted files and enabled near real-time Power BI activity logging via REST APIs.
- Implemented CI/CD to promote Synapse pipelines and stored procedures across environments, parameterizing connections and secrets.
- Defined branching strategies for Synapse and Azure SQL DW so multiple teams could develop in parallel with fewer merge conflicts.

PRIOR EXPERIENCE**ACCENTURE** (May 2017 -June 2020)**Application Development Analyst [Domain: Digital Marketing (US)]**

- Worked as a full-stack developer for a US-based digital marketing company.
- Migrated multiple features from AngularJS to Angular 6 in a team setting.
- Delivered multiple internal applications using Angular, Node.js, and MongoDB.

ACHIEVEMENTS & EXTRA CURRICULAR

- Received the Founder's Award (2023) for exceptional performance while managing multiple clients and projects concurrently.
- Earned multiple spot awards/bonuses for delivery excellence, strong customer outcomes, and end-to-end ownership.
- Placed 3rd in 6th semester and received the Director's Medal for academic excellence.